ANANYA SINGH RAJVANSHI

Uttar Pradesh, India | +91 93360 92755 | singhrajvanshiananya@gmail.com
linkedin.com/in/ananya-singh-rajvanshi-847252276 | github.com/ananya320

PROFESSIONAL SUMMARY

AI Engineer with a B.Tech in Computer Science (graduated July 2026) and hands-on experience building RAG pipelines, LLM-integrated applications, and multi-agent workflows using LangChain, LangGraph, CrewAI, ChromaDB, and FastAPI. Experienced in retrieval systems, agentic orchestration, applied LLM evaluation, and containerized deployment with Docker, CI/CD, and AWS. Internship experience across GenAI, data analytics, and software development. Seeking an AI Engineer role to apply practical LLM integration and applied ML experience to production systems.

TECHNICAL SKILLS

AI/ML & GenAI: LLMs, Retrieval-Augmented Generation (RAG), Prompt Engineering, Embeddings, Vector Search, TF-IDF, Multi-Agent Systems, Agentic Orchestration, Retrieval Evaluation, Bayesian Knowledge Tracing, Expectation-Maximization, NLP, Machine Learning, Deep Learning
Frameworks, APIs & Platforms: LangChain, LangGraph, CrewAI, PyTorch, TensorFlow, Google Gemini API, Gemma 3, Anthropic Claude API, Google Vertex AI, Pandas, NumPy
Programming Languages: Python, Java, C, SQL
Backend, DevOps & Databases: FastAPI, REST APIs, React.js, MongoDB, ChromaDB, Docker, CI/CD, GitHub Actions, Git, AWS (EC2), Render, Vercel, JDBC, Jinja2
Data & Visualization: Power BI, Tableau, Excel, Statistical Analysis

EXPERIENCE

AI Automation Engineer Intern — PerformDigital Sep 2026 – Oct 2026

- Building AI-driven automation workflows using LLMs, APIs, and Gemma 3
- Developing agentic AI integrations to automate business processes and reduce manual effort

RAG Developer — FrugalX Dec 2025 – Sep 2026

- Built a production RAG pipeline for an AI book-writing application, using LangChain and LangGraph for multi-step agentic orchestration, with document chunking, embeddings, and vector-based retrieval via ChromaDB
- Integrated the Gemini API with prompt engineering to reduce hallucination rate by ~10%, achieving 85% relevant chunk retrieval within the top-5 results
- Cut prompt token usage by ~30% by limiting context to the top-3 relevant chunks, lowering inference cost while maintaining output quality
- Containerized the pipeline with Docker and deployed to AWS EC2 through an automated CI/CD workflow, achieving consistent response times of 1–3 seconds

GenAI Intern — Smart Interz (Google Cloud-run Program) Jun 2025 – Nov 2025

- Built GenAI application modules using the Gemini API and ChromaDB for vector storage and retrieval, applying responsible and ethical AI practices throughout
- Designed prompt engineering strategies and AI application architectures through structured, project-based development
- Applied generative AI and NLP techniques to real-world use cases, including automated content generation and workflow automation

Data Analytics Intern — Deloitte Jan 2024 – Apr 2024

- Analyzed datasets exceeding 1,000 records using Python and statistical analysis to identify key trends and derive actionable insights
- Built interactive Tableau dashboards to visualize complex data patterns and support stakeholder decision-making
- Automated data cleaning and preprocessing workflows in Python, improving analysis efficiency

PROJECTS

AI Book Generating App
FastAPI, Google Gemini, LangChain, LangGraph, CI/CD | github.com/ananya320/bookai

- Built a full-stack AI story generator (FastAPI + Google Gemini) using a 5-stage multi-agent pipeline for content generation, editing, and summarization
- Designed a custom CSS 3D interface and persistent story archive with continue/edit/regenerate features
- Deployed on Render with GitHub CI/CD; resolved production issues including SDK migration and secret management

Personalized Learning AI
FastAPI, React.js, MongoDB, TF-IDF, Bayesian Methods | github.com/ananya320/personalized-learning-ai

- Built a full-stack adaptive learning platform (FastAPI/React/MongoDB) using Elo rating and Bayesian Knowledge Tracing to personalize content difficulty and track mastery in real time
- Implemented Expectation-Maximization from scratch to fit ML model parameters from real user data instead of fixed constants
- Built a RAG pipeline (TF-IDF retrieval) enabling AI-generated content grounded in user-uploaded documents
- Deployed a production full-stack app across three cloud platforms (Render, Vercel, MongoDB Atlas), independently debugging real deployment and dependency issues

Additional: AR Shopping Website *(WebAR, Three.js, Blender)* — real-time 3D product visualization with sub-2-second asset load times and custom Blender-created assets. *In progress.*

EDUCATION

B.Tech, Computer Science 2022 – July 2026

Dr. A. P. J. Abdul Kalam Technical University, Lucknow — First Division
Relevant Coursework: Machine Learning, Natural Language Processing, Data Structures & Algorithms, Database Management Systems, Python, React.js

CERTIFICATIONS

- Generative AI Development Program — Smart Interz, run on Google Cloud (2025)
- Anthropic Claude with Google Vertex AI (2026)